

A novel hybrid approach for positive-valued DAG learning

Yao Zhao

YAOZHAO16@TEMPLE.EDU

*Department of Statistics, Operations, and Data Science
Fox School of Business, Temple University
Philadelphia, PA, USA*

Editors: Bijan Mazaheri and Niels Richard Hansen

Abstract

Causal discovery from observational data remains a fundamental challenge in machine learning and statistics, particularly when variables represent inherently positive quantities such as gene expression levels, asset prices, company revenues, or population counts, which naturally follow multiplicative rather than additive dynamics. We propose the Hybrid Moment-Ratio Scoring (H-MRS) algorithm, a novel approach for learning directed acyclic graphs (DAGs) from positive-valued data that combines moment-based scoring with log-scale regression. The key insight is that for positive-valued variables, the moment ratio $\frac{\mathbb{E}[X_j^2]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]}$ provides an effective criterion for causal ordering, where S denotes candidate parent sets. H-MRS combines log-scale Ridge regression for moment-ratio estimation with greedy ordering construction based on raw-scale moment ratios, followed by ElasticNet-based parent selection to recover the final DAG structure. We evaluate H-MRS on synthetic log-linear data, showing that it achieves competitive precision and recall. The algorithm is computationally efficient and naturally respects positivity constraints, making it well-suited for applications such as genomics and economics. Our results highlight that combining log-scale modeling with raw-scale moment ratios offers a practical and robust framework for causal discovery in positive-valued domains.

Keywords: Causal discovery, Directed acyclic graphs, Log-linear models, Moment-ratio scoring, Positive-valued data

1. INTRODUCTION

Causal discovery from observational data remains one of the most fundamental challenges in machine learning and statistics, with applications spanning genomics, economics, epidemiology, and social sciences. The task of learning directed acyclic graphs (DAGs) from strictly positive-valued variables, as commonly occurs in biological expression data, economic indicators, demographic statistics, and social network metrics, presents both a challenge and an opportunity: while standard additive-noise methods are often misspecified in such settings, positivity under a log-linear model can also yield full DAG identifiability rather than only a Markov equivalence class.

The key insight driving our work is that for positive valued variables that exhibit multiplicative relationships, which are common in biological, economic, and financial systems, existing causal discovery methods based on additive noise models are theoretically misspecified. Consider gene regulatory networks where expression of gene j is proportional to the product of regulator concentrations, or asset pricing where returns compound multiplicatively. In such settings, the natural structural equation model is:

$$X_j = \exp \left(\theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k + \epsilon_j \right) \quad (1)$$

rather than the additive form $X_j = \theta_j + \sum_k \beta_{kj} X_k + \epsilon_j$ assumed by most existing methods. Our moment-ratio criterion is specifically designed to exploit the structure of this log-linear model, providing identifiability guarantees that do not hold under model misspecification.

We propose the Hybrid Moment-Ratio Scoring (H-MRS) algorithm, a novel framework that combines log-scale Ridge regression for moment-ratio computation with ElasticNet for parent selection to learn DAGs from positive-valued observational data. Specifically, for node j and candidate set S , we consider the score

$$\mathcal{M}(j, S) = \frac{\mathbb{E}[X_j^2]}{\mathbb{E}[(\mathbb{E}[X_j | S])^2]},$$

which compares an unconditional second moment with a conditional second moment. The conditioning set S appears in the denominator through the conditional expectation. The key property is that this score is minimized when S contains all true parents of node j , as established in Proposition 1. These estimates enable raw-scale moment-ratio scoring to iteratively construct a causal ordering. The resulting criterion is invariant to multiplicative rescaling of individual variables, which helps distinguish it from ordering heuristics driven purely by marginal variance.

The remainder of this paper is organized as follows. Section 2 reviews related work on causal discovery, contrasting score-based, and constraint-based approaches, with particular attention to methods designed for non-Gaussian and positive-valued data. Section 3 presents the H-MRS algorithm in detail, including the log-linear structural equation model, log-scale Ridge regression for conditional expectation estimation, raw-scale moment-ratio scoring for causal ordering, and ElasticNet-based parent selection. Section 4 establishes theoretical properties of the moment-ratio criterion and analyzes the computational complexity of H-MRS. Section 5 evaluates H-MRS on synthetic log-linear data across varying problem scales and sparsity levels. Section 6 applies the method to financial data comprising 2,223 companies and 19 variables, revealing interpretable causal structures in corporate finance. Section 7 concludes with a discussion of limitations and future directions.

2. RELATED WORK

Our work builds upon several interconnected research streams in causal discovery, with particular relevance to methods designed for non-Gaussian and positive-valued data.

Classical Causal Discovery Methods. Traditional causal structure learning approaches fall into two main categories. Constraint-based methods such as PC (Spirites et al., 2000) and its variants rely on conditional independence testing, but standard tests assume Gaussian distributions or use rank-based nonparametric tests that have limited power for detecting multiplicative dependencies in log-linear models. Score-based methods like GES (Chickering, 2002) optimize scoring functions (e.g., BIC, BGe) over graph space, but these scores are derived under additive Gaussian noise assumptions, making them theoretically misspecified when the true data-generating process is multiplicative. While several causal discovery approaches have been developed for nonlinear or non-Gaussian settings, these methods are often computationally intensive and tend to scale poorly to moderate or large graphs.

A major breakthrough came with the Linear Non-Gaussian Acyclic Model (LiNGAM) (Shimizu et al., 2006), which showed that non-Gaussian additive noise enables full causal identifiability beyond the Markov equivalence class. Subsequent work proposed more practical algorithms for estimating such models including the DirectLiNGAM algorithm (Shimizu et al., 2011). However, LiNGAM-type models assume $X_j = \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k + \epsilon_j$, which is fundamentally misspecified for positive-valued data exhibiting multiplicative relationships (e.g., gene expression cascades, compounding financial returns). Our log-linear model $X_j = \exp(\sum_k \beta_{kj} X_k + \epsilon_j)$ naturally captures multiplicative effects while preserving positivity, and our moment-ratio criterion exploits this exponential structure to provide identifiability guarantees that do not hold under additive models.

Moment-based Approaches for Positive Data. Most directly related to our work is the Moments Ratio Scoring (MRS) algorithm developed by Park and Park (2019), which specifically targets high-dimensional Poisson structural equation models. Their approach demonstrates the power of moment-based scoring for discrete count data, achieving polynomial-time recovery with sample complexity $O(d^2 \log^9 p)$. Like our method, MRS decouples ordering estimation from parent search using ℓ_1 -regularized regression and employs moment ratios for scoring.

However, their moment ratio formula is specifically designed for Poisson count data, exploiting the mean-variance equality property of Poisson distributions. In contrast, our H-MRS algorithm uses the simpler ratio $\frac{\mathbb{E}[X_j^2]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]}$ for continuous positive-valued data. Both methods share the key insight that moment ratios are *minimized* when the conditioning set contains the true parents, enabling greedy ordering construction. Our approach combines this principle with log-scale regression for robust conditional expectation estimation in continuous log-linear models.

Post-Nonlinear Causal Models. Our log-linear model (2) is a special case of the post-nonlinear (PNL) framework $x_i = f_{i,2}(f_{i,1}(pa_i) + e_i)$ (Zhang and Hyvarinen, 2012), where we fix $f_{i,2}(\cdot) = \exp(\cdot)$ and constrain $f_{i,1}$ to be linear. While they estimate both functions nonparametrically via mutual information minimization, we exploit the known exponential structure for computational efficiency and develop moment-ratio scoring that leverages the plateau property specific to log-linear models, complementing their general results.

Log-linear and Multiplicative Models. Log-linear models have long been recognized as natural frameworks for positive data, appearing in econometrics (Cameron and Trivedi, 2013), epidemiology, and genomics. However, most existing causal discovery methods do not explicitly leverage the log-linear structure for improved structure learning, and practical algorithms for causal discovery in positive-valued domains remain limited.

Key Distinctions of H-MRS. Our approach differs in several ways: (i) unlike (Park and Park, 2019), we target continuous rather than count data; (ii) we combine log-scale regression with raw-scale moment ratios to capture multiplicative relationships; (iii) positivity is respected natively without strong distributional assumptions.

3. METHODOLOGY

3.1. Problem Formulation

We consider the problem of learning directed acyclic graphs (DAGs) from observational data where all variables are strictly positive-valued. Let $\mathbf{X} = (X_1, \dots, X_p)$ denote a random vector of p positive-valued variables following a joint distribution $P(\mathbf{X})$. Our goal is to recover the causal

DAG $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ where $\mathcal{V} = \{1, \dots, p\}$ represents nodes (variables) and $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$ represents directed edges encoding causal relationships.

For each variable X_j , let $\text{Pa}(j)$ denote its parent set in the true DAG. We assume the data follows a log-linear structural equation model:

$$\log X_j = \theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k + \epsilon_j, \quad j = 1, \dots, p \quad (2)$$

where θ_j is an intercept term, β_{kj} represents the causal effect of X_k on X_j , and ϵ_j is independent noise satisfying the following assumption.

Assumption 1 (Bounded Noise). For each variable X_j , the noise term ϵ_j satisfies:

(A1) Zero mean: $\mathbb{E}[\epsilon_j] = 0$

(A2) Finite variance: $\text{Var}(\epsilon_j) = \sigma_j^2 > 0$

(A3) Almost surely bounded: $|\epsilon_j| \leq B$ for some constant $B > 0$

(A4) Independence: $\epsilon_j \perp\!\!\!\perp X_k$ for all $k \in \text{Pa}(j)$

The boundedness condition (A3) ensures that all moments of X_j exist and are finite, which is essential for the theoretical guarantees in Section 4. Unlike LiNGAM-type approaches, the identifiability of H-MRS does not rely on non-Gaussian noise but instead follows from the log-linear structural form and the moment-ratio plateau property in Proposition 1.

Our log-linear structural equation model (2) employs a specific parametric form that differs from the classical log-log models prevalent in econometrics. We clarify this choice and its implications in Appendix A, where we also provide additional discussion and examples illustrating the modeling assumptions.

3.2. Log-Scale Conditional Expectation Estimation via Ridge Regression

H-MRS first employs Ridge regression on the log-transformed data to approximate conditional expectations for moment-ratio computation. For each variable X_j and candidate parent set $S \subseteq \{1, \dots, p\} \setminus \{j\}$, we estimate:

$$\log X_j \approx \hat{\theta}_j + \sum_{k \in S} \hat{\beta}_{kj} X_k \quad (3)$$

The Ridge estimator provides stable parameter estimates by adding ℓ_2 regularization:

$$\begin{aligned} (\hat{\theta}_j, \hat{\beta}_j) = \arg \min_{\theta, \beta} \frac{1}{2n} \sum_{i=1}^n \left(\log X_j^{(i)} - \theta - \sum_{k \in S} \beta_k X_k^{(i)} \right)^2 \\ + \lambda_{\text{ridge}} \|\beta\|_2^2 \end{aligned} \quad (4)$$

where $\lambda_{\text{ridge}} > 0$ controls the regularization strength and $X_k^{(i)}$ denotes the i -th observation of variable X_k .

The log-scale regression serves two purposes: (1) it captures the multiplicative relationships inherent in positive data through the exponential transformation, and (2) it provides numerical stability by avoiding potential overflow issues when working directly with positive-valued data that

may span several orders of magnitude. For numerical stability, we apply clipping: $\log X_j^{(i)} = \log(\max(X_j^{(i)}, 10^{-10}))$ to avoid logarithms of zero.

From the fitted model, we compute the predicted conditional expectation:

$$\hat{\mu}_{j|S}^{(i)} = \exp \left(\hat{\theta}_j + \sum_{k \in S} \hat{\beta}_{kj} X_k^{(i)} \right) \quad (5)$$

3.3. Raw-Scale Moment-Ratio Scoring for Causal Ordering

H-MRS then computes moment ratios on the original (raw) scale to score causal relationships. For variable X_j with candidate parent set S , we define the moment-ratio score:

$$\mathcal{M}(j, S) = \frac{\mathbb{E}[X_j^2]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]} \quad (6)$$

This criterion satisfies a monotonicity property that enables greedy causal ordering: $\mathcal{M}(j, S)$ is minimized when S contains all parents of j , and achieves the same minimum value for any such superset (the "plateau property" established formally in Proposition 1, Section 4). This enables our greedy selection strategy. At each step, we select the variable whose moment ratio is smallest when conditioning on all previously ordered variables, indicating that its parents are already in the ordering.

In practice, we estimate this score using the outputs from Equation (5):

$$\hat{\mathcal{M}}(j, S) = \frac{\frac{1}{n} \sum_{i=1}^n (X_j^{(i)})^2}{\frac{1}{n} \sum_{i=1}^n (\hat{\mu}_{j|S}^{(i)})^2} \quad (7)$$

The key insight is that while we estimate conditional expectations on the log scale for numerical stability and to capture multiplicative relationships, we compute moment ratios on the raw scale to preserve the theoretical properties necessary for causal ordering. The log-scale regression provides accurate predictions $\hat{\mu}_{j|S}^{(i)}$ which are then used to evaluate the raw-scale moment ratio criterion.

3.4. Parent Selection via Sparsity-Inducing Regression

After establishing a causal ordering through moment-ratio scoring, we select parents for each variable using ElasticNet regression on the log-transformed data. This two-stage design, using Ridge for estimating moment ratios and ElasticNet for selecting parents, addresses distinct statistical requirements.

Why Ridge for moment-ratio computation? Accurate moment-ratio estimation requires *unbiased* conditional expectation estimates to preserve the theoretical ordering guarantees of Proposition 1. Ridge regression provides stable, low-variance predictions without aggressive shrinkage, ensuring that $\hat{\mu}_{j|S}$ faithfully approximates $\mathbb{E}[X_j|S]$ across all candidate sets. The modest ℓ_2 penalty (λ_{ridge}) controls multicollinearity without introducing the selection bias inherent in ℓ_1 penalties.

Why ElasticNet for parent selection? Proposition 1(iii) establishes that the moment ratio achieves its minimum value for *any* superset $S \supseteq \text{Pa}(j)$ that excludes descendants. This plateau property means that moment ratios alone cannot distinguish the true parent set from larger supersets. ElasticNet addresses this limitation by combining ℓ_1 -induced sparsity with an ℓ_2 penalty that

stabilizes selection in the presence of correlated predictors, which are common in applications such as financial and genomic data. In contrast, Lasso tends to select a single variable from a correlated group, which may lead to missing true parents.

Why not use the same regression for both? Using ElasticNet for moment-ratio computation would introduce selection bias because different candidate sets S would shrink different variables to zero, which distorts the moment-ratio comparisons that Proposition 1 relies upon. Conversely, Ridge regression lacks the sparsity needed to identify minimal parent sets from the plateau of equally-scoring supersets. The two-stage approach combines Ridge’s stable prediction for robust ordering with ElasticNet’s variable selection for parsimonious edge recovery.

3.5. Complete H-MRS Algorithm

The H-MRS algorithm exploits the moment-ratio plateau property (Proposition 1) to construct a causal ordering through greedy selection. At each iteration, the variable with minimum moment ratio when conditioned on all previously ordered variables is selected next, as this indicates its true parents are already in the ordering.

Algorithm 1 presents the complete H-MRS procedure. The algorithm alternates between moment-ratio scoring for ordering and ElasticNet regression for parent selection, maintaining the separation between structure learning and parameter estimation. For the first variable in the ordering, we set $S(j) = \emptyset$ for all j . Since no regression is possible without predictors, we compute the moment ratio using the unconditional second moment: $\hat{\mathcal{M}}(j, \emptyset) = \frac{\frac{1}{n} \sum_i (X_j^{(i)})^2}{\left(\frac{1}{n} \sum_i X_j^{(i)}\right)^2}$. The variable with the minimum score is selected as the first in the ordering. This corresponds to selecting the variable with the smallest coefficient of variation, which serves as a natural starting point for the greedy construction.

4. THEORETICAL PROPERTIES

This section establishes the theoretical foundations of H-MRS, providing formal guarantees for the moment-ratio criterion and analyzing the algorithm’s computational complexity.

4.1. Moment-Ratio Identifiability

The correctness of H-MRS relies on the fundamental property that moment ratios can distinguish between correct and incorrect parent sets. We formalize this property under the log-linear structural equation model.

Proposition 1 (Moment-Ratio Plateau Property). Under the log-linear structural equation model (2) with bounded independent noise terms satisfying Assumption 1, for any variable X_j , the moment ratio $\mathcal{M}(j, S)$ satisfies:

- (i) $\mathcal{M}(j, S) \geq 1$ for all $S \subseteq \{1, \dots, p\} \setminus \{j\}$.
 - (ii) For $S_1 \subseteq S_2$: $\mathcal{M}(j, S_2) \leq \mathcal{M}(j, S_1)$ (conditioning on more variables yields smaller or equal moment ratios).
 - (iii) $\mathcal{M}(j, S)$ achieves its minimum value when and only when $\text{Pa}(j) \subseteq S \subseteq \text{NonDesc}(j)$.
- Equivalently:

- If $\text{Pa}(j) \subseteq S \subseteq \text{NonDesc}(j)$, then $\mathcal{M}(j, S) = \mathcal{M}(j, \text{Pa}(j))$ (plateau property)
- If $\text{Pa}(j) \not\subseteq S$, then $\mathcal{M}(j, S) > \mathcal{M}(j, \text{Pa}(j))$ (strict inequality)

Algorithm 1 Hybrid Moment-Ratio Scoring (H-MRS)

Require: Data matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$, Ridge parameter λ_{ridge} , ElasticNet parameters λ, ρ , threshold τ , max degree d_{max}

Ensure: Causal ordering π and parent sets $\{\hat{\text{Pa}}(j)\}_{j=1}^p$

- 1: Initialize $\pi = []$, remaining variables $\mathcal{R} = \{1, \dots, p\}$
- 2: **for** $m = 1$ to p **do**
- 3: **for** each $j \in \mathcal{R}$ **do**
- 4: **if** $m = 1$ **then**
- 5: $S(j) = \emptyset$ {First variable has no parents}
- 6: **else**
- 7: $S(j) = \pi$ {All j use same candidate set}
- 8: **end if**
- 9: Fit Ridge: $\log X_j \sim \sum_{k \in S(j)} \beta_k X_k$
- 10: Compute predictions: $\hat{\mu}_{j|S(j)}^{(i)} = \exp(\hat{\theta}_j + \sum_{k \in S(j)} \hat{\beta}_k X_k^{(i)})$
- 11: Calculate score: $\hat{\mathcal{M}}(j, S(j)) = \frac{\sum_i (X_j^{(i)})^2}{\sum_i (\hat{\mu}_{j|S(j)}^{(i)})^2}$
- 12: **end for**
- 13: $\pi_m = \arg \min_{j \in \mathcal{R}} \hat{\mathcal{M}}(j, S(j))$ {Select variable with minimum score; its parents are already ordered}
- 14: $\pi \leftarrow \pi \cup \{\pi_m\}$, $\mathcal{R} \leftarrow \mathcal{R} \setminus \{\pi_m\}$
- 15: **if** $m > 1$ **then**
- 16: Fit ElasticNet: $\log X_{\pi_m} \sim \sum_{k \in \pi \setminus \{\pi_m\}} \beta_k X_k$
- 17: Threshold: $\mathcal{C} = \{k \in \pi \setminus \{\pi_m\} : |\hat{\beta}_k| > \tau\}$
- 18: Keep top- d_{max} by $|\hat{\beta}_k|$: $\hat{\text{Pa}}(\pi_m) = \text{top}(\mathcal{C}, d_{\text{max}})$
- 19: **else**
- 20: $\hat{\text{Pa}}(\pi_m) = \emptyset$ {First node has no parents}
- 21: **end if**
- 22: **end for**
- 23: **return** $\pi, \{\hat{\text{Pa}}(j)\}_{j=1}^p$

Proof. Please see the appendix B.

4.2. Finite Sample Analysis

We analyze the finite sample behavior of the empirical moment-ratio estimator $\hat{\mathcal{M}}(j, S)$ defined in (7).

The theoretical correctness of H-MRS relies on moment ratios being minimized at true parent sets (Proposition 1). However, we only observe empirical moment ratios computed from finite samples. Proposition 2 establishes that our empirical estimates concentrate around their population values.

Assumption 2 (Regularity for finite-sample analysis). Fix a node j and a candidate set $S \subseteq \{1, \dots, p\} \setminus \{j\}$. We assume:

- (B1) (**Sub-exponential tails**) The variable X_j and the conditional mean $\mu_{j|S} := \mathbb{E}[X_j \mid S]$ are sub-exponential with parameters (ν_j, b_j) and $(\nu_{j,S}, b_{j,S})$, i.e., all moments up to order two exist and Bernstein-type concentration inequalities hold for their empirical averages.
- (B2) (**Design regularity**) The Gram matrix $\Sigma_S := \mathbb{E}[X_S X_S^\top]$ has eigenvalues bounded away from 0 and ∞ , and the regressors X_S are sub-Gaussian.

Proposition 2 (Concentration of empirical moment ratios). Consider the log-linear structural equation model (2) and fix a node j and candidate set S . Let $\hat{\mathcal{M}}(j, S)$ be the empirical moment-ratio estimator defined in (7), and $\mathcal{M}(j, S)$ its population counterpart in (6). Suppose Assumption 2 holds and the denominator $\mathbb{E}[(\mathbb{E}[X_j \mid S])^2]$ is bounded away from zero. Then there exist constants $C_{j,S}, c > 0$ such that, for all $\delta \in (0, 1)$ and all sufficiently large n ,

$$\mathbb{P}\left(|\hat{\mathcal{M}}(j, S) - \mathcal{M}(j, S)| \leq C_{j,S} \sqrt{\frac{\log(1/\delta)}{n}}\right) \geq 1 - \delta. \quad (8)$$

In particular, $\hat{\mathcal{M}}(j, S)$ converges to $\mathcal{M}(j, S)$ at the usual parametric rate $O_p(n^{-1/2})$.

Proof. Please see the appendix C.

For H-MRS to correctly order variables, the empirical moment ratios must be accurate enough to distinguish between correct parent sets (which minimize $\mathcal{M}$) and incorrect ones (which have strictly larger $\mathcal{M}$). Proposition 2 guarantees this happens with high probability when the sample size is sufficiently large relative to the separation between correct and incorrect scores.

4.3. Computational Complexity Analysis

We analyze the computational requirements of H-MRS compared to existing DAG learning methods.

Proposition 3 (Time Complexity). The time complexity of H-MRS is $O(p^2 \cdot T_{\text{Ridge}} + p \cdot T_{\text{ElasticNet}})$ where $T_{\text{Ridge}} = O(nq^2 + q^3)$ represents the cost of fitting Ridge regression and $T_{\text{ElasticNet}} = O(nq^2 \cdot K)$ represents the cost of fitting ElasticNet with q predictors over K coordinate descent iterations.

Proof. Please see the appendix D.

Space Complexity. H-MRS requires $O(np + p^2)$ space: $O(np)$ for the data matrix and $O(p^2)$ for storing the estimated adjacency matrix and intermediate regression coefficients.

At the end of this section, we give the remark on hyperparameter requirements. For asymptotic consistency of H-MRS, the regularization parameters should satisfy $\lambda_{\text{ridge}}, \lambda_n \rightarrow 0$ and $\lambda_{\text{ridge}} \sqrt{n}, \lambda_n \sqrt{n} \rightarrow \infty$ as $n \rightarrow \infty$, ensuring Ridge and ElasticNet estimators converge while maintaining stability. The threshold parameter should satisfy $\tau \gtrsim \sqrt{\frac{\log p}{n}}$ for correct edge selection. See Appendix E for detailed hyperparameter guidance and practical selection strategies.

5. EXPERIMENTS

We evaluate H-MRS on synthetic log-linear data to assess its performance in recovering causal structures from positive-valued observations. Our experiments demonstrate that H-MRS effectively recovers causal structures from positive-valued data, highlighting its robustness and practical utility in the target domain.

5.1. Experimental Setup

Data Generation. We generate synthetic data following the log-linear structural equation model (2). For each experiment, we first sample a random topological ordering of p variables, then construct a DAG by randomly adding edges from earlier to later variables in the ordering. For the simulation study, the maximum in-degree d is specified as part of the data-generating process to control structural complexity, with $d \in \{1, 2\}$ corresponding to the simple and complex regimes reported below.

The data generation process follows these steps:

1. Sample intercepts $\theta_j \sim \text{Uniform}(0.5, 2.0)$
2. Sample edge weights $\beta_{kj} \sim \text{Uniform}(-0.3, 0.3)$
3. For each variable j in topological order, generate observations:

$$X_j^{(i)} = \exp \left(\theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k^{(i)} + \epsilon_j^{(i)} \right) \quad (9)$$

where $\epsilon_j^{(i)} \sim \text{Uniform}(-B, B)$ with $B = 0.5$.

This ensures: (i) all variables are strictly positive, (ii) all moments exist and are finite, and (iii) the log-linear relationships that H-MRS exploits are preserved.

Remark on scaling and varsortability. One may ask whether the strong ordering performance could be driven by marginal variance differences rather than the proposed criterion itself, as in the varsortability phenomenon. In our setting, the moment-ratio score

$$\mathcal{M}(j, S) = \frac{\mathbb{E}[X_j^2]}{\mathbb{E}[(\mathbb{E}[X_j | S])^2]}$$

is scale-invariant: if X_j is replaced by cX_j for any constant $c > 0$, then both the numerator and denominator are multiplied by c^2 , leaving the score unchanged. Thus, the ordering criterion is not determined by absolute variable scale alone. Moreover, the simulation model includes both positive and negative edge weights, $\beta_{kj} \sim \text{Uniform}(-0.3, 0.3)$, so child variables do not necessarily exhibit larger marginal variance than their parents. This reduces the possibility that correct ordering is recovered merely from a monotone variance pattern along the causal graph.

Evaluation Metrics. We use standard graph recovery metrics: (1) **Structural Hamming Distance (SHD)**: the number of edge additions, deletions, or reversals required to transform the estimated graph into the true graph; (2) **Precision**: the fraction of estimated edges that are correct; (3) **Recall**: the fraction of true edges correctly identified; (4) **F1-score**: the harmonic mean of precision and recall.

Note that PC and GES return a completed partially directed acyclic graph (CPDAG) rather than a fully directed DAG. To enable comparison with methods that output DAGs, we represent undirected edges in the CPDAG as bidirected edges in the adjacency matrix. This treatment preserves adjacency information when computing recall, while appropriately penalizing precision and SHD due to the absence of a uniquely determined edge direction.

5.2. Performance Analysis

We systematically evaluate H-MRS across different experimental conditions with comprehensive hyperparameter optimization. Our evaluation covers graph sizes $p \in \{10, 20, 30\}$ and sample sizes $n = 500$.

We evaluate H-MRS against three established causal discovery algorithms on log-linear data: PC (constraint-based), GES (scorebased), and DirectLiNGAM (designed for non-Gaussian linear models).¹ These baseline methods are applied using their standard implementations and default conditional independence tests or scoring functions. Specifically, we used the `causal-learn` library (Python interface to the Tetrad project): PC was run with its default conditional independence test and significance level, GES with its default score-based configuration, and DirectLiNGAM using its standard implementation following (Shimizu et al., 2011), without additional hyperparameter tuning. While PC and GES can in principle be combined with alternative tests or scores tailored to nonlinear or non-Gaussian settings, our goal is to compare H-MRS against widely used causal discovery baselines as they are typically applied in practice. Table 1 and 2 present the optimized performance across all tested configurations. H-MRS demonstrates strong performance with F1-scores ranging from 0.733 to 0.800 and correspondingly low SHD values. The algorithm maintains high precision (0.750-1.000) across all settings while achieving reasonable recall, indicating effective control of false positive rates.

Table 1: Simple structures with maximum in-degree $d = 1$

Configuration	Method	SHD	Precision	Recall	F1
p=10, d=1	H-MRS	1	1.000	0.667	0.800
	PC	5	0.375	1.000	0.545
	GES	14	0.143	0.667	0.235
	DirectLiNGAM	5	0.714	0.625	0.667
p=20, d=1	H-MRS	5	0.750	0.750	0.750
	PC	22	0.27	1.00	0.42
	GES	31	0.194	0.875	0.318
	DirectLiNGAM	48	0.380	0.528	0.442
p=30, d=1	H-MRS	10	0.846	0.647	0.733
	PC	58	0.15	0.83	0.26
	GES	76	0.093	0.583	0.161
	DirectLiNGAM	45	0.356	0.500	0.416

The results reveal several important patterns. For simple structures ($d=1$), H-MRS maintains strong performance across scales, with F1-scores declining modestly from 0.800 to 0.733 as problem size increases from 10 to 30 variables. For complex structures ($d=2$), H-MRS demonstrates even stronger performance, achieving F1-scores ranging from 0.745 to 0.900. The algorithm maintains perfect precision (1.000) on smaller problems ($p=10, 20$) while achieving high recall, and sustains competitive performance ($F1=0.745$) even on dense 30-variable networks.

1. PC and GES are included as standard baseline methods in causal discovery, although their classical formulations assume linear Gaussian models and are therefore misspecified for log-linear multiplicative data.

Table 2: Complex structures with maximum in-degree $d = 2$

Configuration	Method	SHD	Precision	Recall	F1
p=10, d=2	H-MRS	2	1.000	0.818	0.900
	PC	13	0.45	0.818	0.581
	GES	12	0.600	0.545	0.571
	DirectLiNGAM	13	0.250	0.429	0.316
p=20, d=2	H-MRS	3	1.000	0.625	0.769
	PC	28	0.222	1.000	0.364
	GES	17	0.308	0.500	0.381
	DirectLiNGAM	63	0.293	0.436	0.351
p=30, d=2	H-MRS	15	0.760	0.731	0.745
	PC	47	0.319	0.720	0.439
	GES	38	0.423	0.498	0.456
	DirectLiNGAM	69	0.328	0.417	0.367

6. REAL DATA ANALYSIS

6.1. Data Description

We apply H-MRS to real-world financial data to evaluate its performance on genuine observational data from economic systems. The dataset comprises balance sheet, income statement, and cash flow information for publicly traded companies, sourced from a publicly available financial dataset on Kaggle (<https://www.kaggle.com/datasets/pacificrm/financial-sheets>).

Starting with 4,668 companies and multiple financial tables, we perform the following preprocessing steps: (1) merge balance sheet, income statement, and cash flow data by company identifier; (2) select 19 key financial variables spanning assets, liabilities, equity, income, expenses, and market valuation metrics; (3) remove companies with missing values. After preprocessing, we obtain a clean cross-sectional dataset with $n = 2,223$ companies and $p = 19$ variables.

The 19 selected variables capture comprehensive financial information across multiple categories:

Assets (6 variables): Total Assets, Current Assets, Net Fixed Assets, Inventory, Accounts Receivable, Cash and Equivalents

Liabilities & Equity (5 variables): Total Debt, Current Liabilities, Equity Capital, Reserves, Accounts Payable

Income & Expenses (6 variables): Revenue/Sales, Operating Profit, EBIT, Net Profit, Interest Expense, Depreciation

Market Valuation (2 variables): Market Capitalization, Enterprise Value

All variables are measured in monetary units (millions) and are strictly positive, making the dataset well-suited for the log-linear structural equation model underlying H-MRS.

6.2. Results

The estimated DAG is shown in Appendix F (Figure 1). A salient pattern in the estimated graph is the central role of Equity Capital as an upstream driver. Equity appears as the unique source node and exerts broad influence on operating and valuation variables, including EBIT, Operating Profit,

Inventory, Market Capitalization, Enterprise Value, and Total Assets. This structure is consistent with the economic interpretation that the financing base determines the scale of operations (Fazzari et al., 1987), which subsequently propagates into both accounting aggregates and market valuation. In this sense, equity serves as the foundational resource from which productive capacity and firm value are built.

A second prominent feature is the pervasive influence of Interest Expense, which emerges as a system-wide driver with fourteen outgoing edges. Beyond its direct connection to profitability measures (EBIT, Operating Profit), interest expense influences Current Assets, Current Liabilities, Total Debt, Total Assets, and the valuation variables. This suggests that the cost of debt financing functions as a global constraint, shaping liquidity management, leverage, and ultimately market outcomes. The interpretation aligns with the notion that financing costs transmit broadly across the balance sheet, affecting both the asset side and market capitalization.

Further, the relationship between profitability and leverage is reflected in the edge Operating Profit $\rightarrow$ Total Debt, alongside the direct influence of interest expense. This configuration suggests that profitable firms may support higher levels of borrowing, consistent with creditworthiness arguments (Merton, 1974), while interest expense continues to reflect the cost of carrying such debt. Together, these edges illustrate a tension between profitability as a facilitator of borrowing and interest payments as a constraint on debt capacity.

Finally, Depreciation and Net Fixed Assets are both downstream of equity capital and interest expense, underscoring the link between financing availability, capital intensity, and long-lived assets. At the terminal stage of the ordering, Market Capitalization and Enterprise Value are jointly determined by equity capital and interest expense. This structure resonates with financial theory (Myers, 1984): equity availability provides the foundation for valuation, while borrowing costs act as a discounting channel, moderating the translation of accounting fundamentals into market value.

Taken together, the recovered DAG suggests a coherent narrative in which equity capital and interest expense form dual upstream levers of the financial system: equity provides the base for expansion, while interest expense imposes a cost of capital that permeates the system. These factors cascade through working capital, profitability, and asset accumulation, ultimately shaping firm valuation.

7. DISCUSSION AND CONCLUSION

This paper introduced the Hybrid Moment-Ratio Scoring (H-MRS) algorithm, a novel method for causal discovery in positive-valued domains. By integrating log-scale regression with moment-ratio-based ordering and sparse parent selection, H-MRS leverages the multiplicative structure of positive data while preserving the theoretical identifiability properties of moment ratios. The algorithm operates in polynomial time, produces sparse and interpretable graphs, and naturally respects positivity constraints.

Our empirical study demonstrated that H-MRS achieves competitive performance on synthetic log-linear data, with strong precision and recall across a range of graph sizes and sparsity levels. On real financial data consisting of 19 key firm-level variables, the learned structure highlights economically coherent channels: *Equity Capital* emerged as a foundational source influencing profitability, working-capital components, and valuation, while *Interest Expense* acted as a pervasive upstream driver reflecting the system-wide role of financing costs. These findings illustrate that H-MRS is capable of uncovering meaningful and interpretable causal pathways in real-world economic systems.

Several limitations warrant discussion. First, the method is evaluated in cross-sectional settings; temporal extensions could improve causal interpretation in dynamic systems. Second, the current implementation imposes a maximum in-degree constraint to control complexity, which may omit higher-order interactions. Third, the current formulation assumes strictly positive-valued variables. In applications such as genomic read-count data, observations may be zero-inflated due to biological or measurement processes. Extending the moment-ratio framework to accommodate structural zeros, for example through hurdle or zero-inflated models that separately model the occurrence of zeros and the positive component, represents an interesting direction for future research. Fourth, H-MRS assumes an underlying directed acyclic graph (DAG) structure. In domains such as economics and social systems, feedback loops and cyclic relationships are common. When such cycles are present, the DAG assumption is violated, and H-MRS should be interpreted as providing an approximate acyclic representation rather than recovering the true cyclic structure. In practice, the method may still yield useful insights into dominant directional dependencies, but its outputs should be interpreted with caution under model misspecification. Extending moment-ratio approaches to cyclic or equilibrium-based models is an important direction for future work.

From a practical perspective, users may wish to assess whether their data are compatible with the modeling assumptions of H-MRS. Since the method relies on multiplicative relationships on the original scale (equivalently, additive structure on the log scale), a simple diagnostic is to examine whether log-transformed variables exhibit approximately linear conditional relationships and stabilized variance. In practice, fitting linear models on the log scale and inspecting residual patterns can provide a useful heuristic: well-behaved residuals and reduced heteroscedasticity suggest that the log-linear approximation is reasonable. These checks are informal but can help determine whether H-MRS is appropriate for a given dataset.

In conclusion, H-MRS provides a principled and computationally efficient approach to causal discovery from positive-valued data. By bridging log-scale modeling and moment-ratio criteria, it offers a new perspective on learning directed acyclic graphs in settings where positivity and multiplicative interactions are fundamental. We hope this work stimulates further research on moment-based methods for causal inference and broadens the scope of applications in machine learning, statistics, and the applied sciences.

Code availability. The Python implementation of the H-MRS algorithm and the simulation scripts used in this paper are publicly available at

<https://github.com/statsYAO/H-MRS>.

Acknowledgments

The author would like to thank his advisor, Kuang-Yao Lee, for his invaluable guidance and continuous support throughout his Ph.D. studies.

References

- Uri Alon. *An introduction to systems biology: design principles of biological circuits*. Chapman and Hall/CRC, 2019.
- Peter Bühlmann and Sara van de Geer. *Statistics for High-Dimensional Data: Methods, Theory and Applications*. Springer, Berlin, Heidelberg, 2011. doi: 10.1007/978-3-642-20192-9.

- A Colin Cameron and Pravin K Trivedi. *Regression analysis of count data*, volume 53. Cambridge university press, 2013.
- David Maxwell Chickering. Optimal structure identification with greedy search. *Journal of Machine Learning Research*, 3:507–554, 2002.
- Rick Durrett. *Probability: theory and examples*, volume 49. Cambridge university press, 2019.
- Steven Fazzari, R Glenn Hubbard, and Bruce C Petersen. Financing constraints and corporate investment. Technical report, National Bureau of Economic Research, 1987.
- Zhiguo He and Arvind Krishnamurthy. Intermediary asset pricing. *American Economic Review*, 103(2):732–770, 2013.
- Robert Merton. On the pricing of corporate debt. *The Journal of Finance*, 1974.
- Stewart Myers. The capital structure puzzle. *The Journal of Finance*, 1984.
- Gunwoong Park and Seyoung Park. High-dimensional poisson structural equation model learning via ℓ_1 -regularized regression. *Journal of Machine Learning Research*, 20(95):1–41, 2019.
- Shohei Shimizu, Patrik O Hoyer, Aapo Hyvärinen, Antti Kerminen, and Michael Jordan. A linear non-gaussian acyclic model for causal discovery. *Journal of Machine Learning Research*, 7(10), 2006.
- Shohei Shimizu, Takanori Inazumi, Yasuhiro Sogawa, Aapo Hyvärinen, Yoshinobu Kawahara, Takashi Washio, Patrik Hoyer, and Kenneth Bollen. Directlingam: A direct method for learning a linear non-gaussian structural equation model. *Journal of Machine Learning Research*, 12: 1225–1248, 2011.
- Peter Spirtes, Clark N Glymour, and Richard Scheines. *Causation, prediction, and search*. MIT Press, 2nd edition, 2000.
- Kun Zhang and Aapo Hyvarinen. On the identifiability of the post-nonlinear causal model. *arXiv preprint arXiv:1205.2599*, 2012.
- Hui Zou and Trevor Hastie. Regularization and variable selection via the elastic net. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 67(2):301–320, 2005.

Appendix A. Model Specification and Comparison

In econometrics and production theory, ”log-log” typically refers to models where both the response and predictors are logtransformed:

$$\log X_j = \theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} \log X_k + \epsilon_j$$

This yields the multiplicative (Cobb-Douglas) form:

$$X_j = \exp(\theta_j) \cdot \prod_{k \in \text{Pa}(j)} X_k^{\beta_{kj}} \cdot \exp(\epsilon_j)$$

Such models naturally capture elasticities (β_{kj} represents the percentage change in X_j per 1% change in X_k) and are widely used for production functions, demand curves, and growth models where proportional relationships dominate.

In contrast, our semi-log model specifies:

$$\log X_j = \theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k + \epsilon_j$$

equivalently:

$$X_j = \exp \left(\theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k + \epsilon_j \right)$$

Here, $\log X_j$ depends linearly on the raw values X_k , not their logarithms. This specification arises naturally in several important domains:

1. Gene Regulatory Networks: Transcription factor binding often follows Hill kinetics or linear-exponential dose-response relationships, where log-expression of a target gene depends approximately linearly on the concentration (not log-concentration) of regulatory proteins (Alon, 2019).
2. Financial Contagion Models: Log-returns or log-volatility may depend linearly on the raw leverage ratios or balance sheet quantities of connected institutions, capturing threshold effects and cascading failures (He and Krishnamurthy, 2013).

The semi-log specification offers several advantages for causal discovery:

Identifiability: The linear dependence on (not $\log X_k$) creates asymmetries in conditional distributions that enable causal direction identification, similar to the nonGaussian linear case but adapted for positive data.

Bounded Influence: When parent variables X_k are bounded, $\beta_{kj} X_k$ remains in a compact set, ensuring that all moments of X_j exist and are finite. In contrast, log-log models can produce heavy-tailed distributions when $\beta_{kj} < 0$ and X_k is near zero.

Moment-Ratio Properties: The plateau property (Proposition 1) exploits the specific exponential structure of our model. Under log-log specifications, the moment-ratio criterion would require different theoretical analysis and may lose its ordering guarantees.

The semi-log form is particularly suited for settings where:

1. Variables are positive-valued;
2. Causal effects operate through concentration-dependent rather than elasticity-dependent mechanisms;
3. The data exhibit exponential growth or decay with respect to predictor levels.

For applications where elasticity interpretations are paramount (e.g., demand curves, production functions), the log-log specification may be more appropriate. However, for domains like gene regulation, chemical kinetics, or systems with threshold effects, our semi-log model provides a more natural representation of the underlying causal mechanisms.

Appendix B. Proof of Proposition 1

Proof. By the law of total variance:

$$\text{Var}(X_j) = \mathbb{E}[\text{Var}(X_j|S)] + \text{Var}(\mathbb{E}[X_j|S]) \quad (10)$$

Since $\mathbb{E}[X_j^2] = \text{Var}(X_j) + (\mathbb{E}[X_j])^2$, and by the law of total variance $\text{Var}(X_j) = \mathbb{E}[\text{Var}(X_j|S)] + \text{Var}(\mathbb{E}[X_j|S])$:

$$\mathbb{E}[X_j^2] = \mathbb{E}[\text{Var}(X_j|S)] + \text{Var}(\mathbb{E}[X_j|S]) + (\mathbb{E}[X_j])^2 \quad (11)$$

Using the variance decomposition $\text{Var}(Y) = \mathbb{E}[Y^2] - (\mathbb{E}[Y])^2$ with $Y = \mathbb{E}[X_j|S]$:

$$\begin{aligned} \text{Var}(\mathbb{E}[X_j | S]) &= \mathbb{E}[(\mathbb{E}[X_j | S])^2] - (\mathbb{E}[\mathbb{E}[X_j | S]])^2 \\ &= \mathbb{E}[(\mathbb{E}[X_j | S])^2] - (\mathbb{E}[X_j])^2. \end{aligned} \quad (12)$$

where the last equality uses the tower property $\mathbb{E}[\mathbb{E}[X_j|S]] = \mathbb{E}[X_j]$.

Substituting back:

$$\mathbb{E}[X_j^2] = \mathbb{E}[\text{Var}(X_j|S)] + \mathbb{E}[(\mathbb{E}[X_j|S])^2] \quad (13)$$

Therefore:

$$\mathcal{M}(j, S) = \frac{\mathbb{E}[X_j^2]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]} = 1 + \frac{\mathbb{E}[\text{Var}(X_j|S)]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]} \quad (14)$$

(i) Since $\mathbb{E}[\text{Var}(X_j|S)] \geq 0$, we have $\mathcal{M}(j, S) \geq 1$.

(ii) Let $\mathcal{F}_1 = \sigma(S_1)$ and $\mathcal{F}_2 = \sigma(S_2)$ denote the σ -algebras generated by S_1 and S_2 respectively. Since $S_1 \subseteq S_2$, we have $\mathcal{F}_1 \subseteq \mathcal{F}_2$.

By the definition of conditional variance:

$$\text{Var}(X_j | \mathcal{F}) = \mathbb{E}[X_j^2 | \mathcal{F}] - (\mathbb{E}[X_j | \mathcal{F}])^2 \quad (15)$$

Taking expectations on both sides:

$$\begin{aligned} \mathbb{E}[\text{Var}(X_j | \mathcal{F})] &= \mathbb{E}[\mathbb{E}[X_j^2 | \mathcal{F}] - \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}])^2]] \\ &= \mathbb{E}[X_j^2] - \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}])^2] \end{aligned} \quad (16)$$

where the second equality uses the tower property $\mathbb{E}[\mathbb{E}[X_j^2 | \mathcal{F}]] = \mathbb{E}[X_j^2]$.

Applying (16) to both $\mathcal{F}_1$ and $\mathcal{F}_2$:

$$\mathbb{E}[\text{Var}(X_j | \mathcal{F}_k)] = \mathbb{E}[X_j^2] - \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}_k])^2], \quad k = 1, 2 \quad (17)$$

Therefore:

$$\begin{aligned} &\mathbb{E}[\text{Var}(X_j | \mathcal{F}_1)] - \mathbb{E}[\text{Var}(X_j | \mathcal{F}_2)] \\ &= (\mathbb{E}[X_j^2] - \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}_1])^2]) - (\mathbb{E}[X_j^2] - \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}_2])^2]) \\ &= \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}_2])^2] - \mathbb{E}[(\mathbb{E}[X_j | \mathcal{F}_1])^2] \end{aligned} \quad (18)$$

Define $Z := \mathbb{E}[X_j \mid \mathcal{F}_2]$. By Theorem 4.1.13(ii) in [Durrett \(2019\)](#), since $\mathcal{F}_1 \subseteq \mathcal{F}_2$:

$$\mathbb{E}[Z \mid \mathcal{F}_1] = \mathbb{E}[\mathbb{E}[X_j \mid \mathcal{F}_2] \mid \mathcal{F}_1] = \mathbb{E}[X_j \mid \mathcal{F}_1] \quad (19)$$

By Exercise 4.4.5 in [Durrett \(2019\)](#):

$$\mathbb{E}[(Z - \mathbb{E}[Z \mid \mathcal{F}_1])^2] = \mathbb{E}[Z^2] - \mathbb{E}[(\mathbb{E}[Z \mid \mathcal{F}_1])^2] \quad (20)$$

Substituting $Z = \mathbb{E}[X_j \mid \mathcal{F}_2]$ and $\mathbb{E}[Z \mid \mathcal{F}_1] = \mathbb{E}[X_j \mid \mathcal{F}_1]$:

$$\begin{aligned} & \mathbb{E}[(\mathbb{E}[X_j \mid \mathcal{F}_2] - \mathbb{E}[X_j \mid \mathcal{F}_1])^2] \\ &= \mathbb{E}[(\mathbb{E}[X_j \mid \mathcal{F}_2])^2] - \mathbb{E}[(\mathbb{E}[X_j \mid \mathcal{F}_1])^2] \end{aligned} \quad (21)$$

Comparing (18) and (21), we obtain:

$$\begin{aligned} & \mathbb{E}[\text{Var}(X_j \mid \mathcal{F}_1)] - \mathbb{E}[\text{Var}(X_j \mid \mathcal{F}_2)] \\ &= \mathbb{E}[(\mathbb{E}[X_j \mid \mathcal{F}_2] - \mathbb{E}[X_j \mid \mathcal{F}_1])^2] \end{aligned} \quad (22)$$

On the other hand, by the definition of conditional variance:

$$\text{Var}(Z \mid \mathcal{F}_1) = \mathbb{E}[(Z - \mathbb{E}[Z \mid \mathcal{F}_1])^2 \mid \mathcal{F}_1] \quad (23)$$

Taking expectations on both sides:

$$\mathbb{E}[\text{Var}(Z \mid \mathcal{F}_1)] = \mathbb{E}[(Z - \mathbb{E}[Z \mid \mathcal{F}_1])^2] \quad (24)$$

Applying (24) with $Z = \mathbb{E}[X_j \mid \mathcal{F}_2]$ and using $\mathbb{E}[Z \mid \mathcal{F}_1] = \mathbb{E}[X_j \mid \mathcal{F}_1]$:

$$\mathbb{E}[(\mathbb{E}[X_j \mid \mathcal{F}_2] - \mathbb{E}[X_j \mid \mathcal{F}_1])^2] = \mathbb{E}[\text{Var}(\mathbb{E}[X_j \mid \mathcal{F}_2] \mid \mathcal{F}_1)] \quad (25)$$

Combining (22) and (25):

$$\begin{aligned} & \mathbb{E}[\text{Var}(X_j \mid \mathcal{F}_1)] - \mathbb{E}[\text{Var}(X_j \mid \mathcal{F}_2)] \\ &= \mathbb{E}[\text{Var}(\mathbb{E}[X_j \mid \mathcal{F}_2] \mid \mathcal{F}_1)] \end{aligned} \quad (26)$$

Rearranging:

$$\mathbb{E}[\text{Var}(X_j \mid \mathcal{F}_1)] = \mathbb{E}[\text{Var}(X_j \mid \mathcal{F}_2)] + \mathbb{E}[\text{Var}(\mathbb{E}[X_j \mid \mathcal{F}_2] \mid \mathcal{F}_1)] \quad (27)$$

Returning to our original notation with S_1 and S_2 :

$$\begin{aligned} & \mathbb{E}[\text{Var}(X_j \mid S_1)] \\ &= \mathbb{E}[\text{Var}(X_j \mid S_2)] + \mathbb{E}[\text{Var}(\mathbb{E}[X_j \mid S_2] \mid S_1)] \\ &\geq \mathbb{E}[\text{Var}(X_j \mid S_2)] \end{aligned} \quad (28)$$

where the inequality follows since variance is non-negative.

Moreover, by Theorem 4.1.13(ii) in [Durrett \(2019\)](#), $\mathbb{E}[X_j | S_1] = \mathbb{E}[\mathbb{E}[X_j | S_2] | S_1]$. Applying the conditional Jensen inequality to the convex function $\varphi(x) = x^2$:

$$\begin{aligned} & \mathbb{E}[(\mathbb{E}[X_j | S_1])^2] \\ &= \mathbb{E}[(\mathbb{E}[\mathbb{E}[X_j | S_2] | S_1])^2] \\ &\leq \mathbb{E}[\mathbb{E}[(\mathbb{E}[X_j | S_2])^2 | S_1]] \\ &= \mathbb{E}[(\mathbb{E}[X_j | S_2])^2] \end{aligned} \tag{29}$$

where the inequality follows from the conditional Jensen inequality $(\mathbb{E}[Y | \mathcal{F}])^2 \leq \mathbb{E}[Y^2 | \mathcal{F}]$ applied to $Y = \mathbb{E}[X_j | S_2]$ and $\mathcal{F} = \sigma(S_1)$, and the last equality uses the tower property.

Therefore, from the representation in (i), $\mathcal{M}(j, S) = 1 + \frac{\mathbb{E}[\text{Var}(X_j | S)]}{\mathbb{E}[(\mathbb{E}[X_j | S])^2]}$:

$$\begin{aligned} \mathcal{M}(j, S_2) &= 1 + \frac{\mathbb{E}[\text{Var}(X_j | S_2)]}{\mathbb{E}[(\mathbb{E}[X_j | S_2])^2]} \\ &\leq 1 + \frac{\mathbb{E}[\text{Var}(X_j | S_1)]}{\mathbb{E}[(\mathbb{E}[X_j | S_1])^2]} \\ &= \mathcal{M}(j, S_1) \end{aligned} \tag{30}$$

Thus $\mathcal{M}(j, S_2) \leq \mathcal{M}(j, S_1)$.

(iii) By the local Markov property of DAGs, $X_j \perp\!\!\!\perp (\text{NonDesc}(j) \setminus \text{Pa}(j)) | \text{Pa}(j)$.

Therefore, for any S satisfying $\text{Pa}(j) \subseteq S$ and $S \cap \text{Desc}(j) = \emptyset$ (i.e., S contains all parents but no descendants of j), we have:

$$\mathbb{E}[X_j | S] = \mathbb{E}[X_j | \text{Pa}(j)] \tag{31}$$

This follows because $S \setminus \text{Pa}(j) \subseteq \text{NonDesc}(j) \setminus \text{Pa}(j)$ by the condition $S \cap \text{Desc}(j) = \emptyset$, which are conditionally independent of X_j given $\text{Pa}(j)$.

Consequently:

$$\text{Var}(X_j | S) = \text{Var}(X_j | \text{Pa}(j)) > 0 \tag{32}$$

where the strict inequality holds due to the independent noise ϵ_j in model (2).

This shows $\mathcal{M}(j, S) = \mathcal{M}(j, \text{Pa}(j))$ for all S satisfying $\text{Pa}(j) \subseteq S \subseteq \text{NonDesc}(j)$. Combined with (ii), this establishes that the moment ratio achieves its minimum value on this plateau. $\square$

Now I prove the part 2: Strict inequality when parents are missing

Now consider S such that $\text{Pa}(j) \not\subseteq S$. Let $M = \text{Pa}(j) \setminus S \neq \emptyset$ denote the set of missing parents.

Key Lemma: Under the log-linear model, if $M \neq \emptyset$ and all $\beta_{kj} \neq 0$ for $k \in \text{Pa}(j)$, then:

$$\mathbb{E}[\text{Var}(X_j | S)] > \mathbb{E}[\text{Var}(X_j | \text{Pa}(j))] \tag{33}$$

Proof of Lemma:

By the conditional variance decomposition for nested σ -algebras, applied to $\mathcal{F}_1 = \sigma(S)$ and $\mathcal{F}_2 = \sigma(S \cup \text{Pa}(j))$, we have

$$\mathbb{E}[\text{Var}(X_j | S)] = \mathbb{E}[\text{Var}(X_j | S, \text{Pa}(j))] + \mathbb{E}[\text{Var}(\mathbb{E}[X_j | S, \text{Pa}(j)] | S)].$$

Since $S \subseteq \text{NonDesc}(j)$, by the local Markov property

$$X_j \perp\!\!\!\perp (S \setminus \text{Pa}(j)) | \text{Pa}(j)$$

which implies $\mathbb{E}[X_j | S, \text{Pa}(j)] = \mathbb{E}[X_j | \text{Pa}(j)]$ and $\text{Var}(X_j | S, \text{Pa}(j)) = \text{Var}(X_j | \text{Pa}(j))$. Hence

$$\mathbb{E}[\text{Var}(X_j | S)] = \mathbb{E}[\text{Var}(X_j | \text{Pa}(j))] + \mathbb{E}[\text{Var}(\mathbb{E}[X_j | \text{Pa}(j)] | S)].$$

We need to show that the second term is strictly positive when $M \neq \emptyset$.

From the log-linear model:

$$\mathbb{E}[X_j | \text{Pa}(j)] = \mathbb{E} \left[\exp \left(\theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k + \epsilon_j \right) \middle| \text{Pa}(j) \right] \quad (34)$$

$$= \exp \left(\theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k \right) \cdot \mathbb{E}[\exp(\epsilon_j)] \quad (35)$$

Since ϵ_j is independent of all parent variables and has non-degenerate distribution (with $\text{Var}(\epsilon_j) = \sigma_j^2 > 0$), $\mathbb{E}[\exp(\epsilon_j)]$ is a positive constant, say $c_j = \mathbb{E}[\exp(\epsilon_j)] > 0$.

Therefore:

$$\mathbb{E}[X_j | \text{Pa}(j)] = c_j \exp \left(\theta_j + \sum_{k \in \text{Pa}(j)} \beta_{kj} X_k \right) \quad (36)$$

Now, consider $\text{Var}(\mathbb{E}[X_j | \text{Pa}(j)] | S)$.

Case 1: S contains some but not all parents

Without loss of generality, partition $\text{Pa}(j) = S_{\text{in}} \cup M$ where $S_{\text{in}} = S \cap \text{Pa}(j)$ and $M = \text{Pa}(j) \setminus S \neq \emptyset$.

Then:

$$\mathbb{E}[X_j | \text{Pa}(j)] = c_j \exp \left(\theta_j + \sum_{k \in S_{\text{in}}} \beta_{kj} X_k + \sum_{k \in M} \beta_{kj} X_k \right) \quad (37)$$

Conditioning on S :

$$\mathbb{E}[\mathbb{E}[X_j | \text{Pa}(j)] | S] = c_j \exp \left(\theta_j + \sum_{k \in S_{\text{in}}} \beta_{kj} X_k \right) \cdot \mathbb{E} \left[\exp \left(\sum_{k \in M} \beta_{kj} X_k \right) \middle| S \right] \quad (38)$$

The key observation is that $\mathbb{E}[X_j | \text{Pa}(j)]$ depends on the missing parents X_k for $k \in M$ through the exponential function. Since the exponential is strictly convex and the X_k have non-degenerate conditional distributions given S (they are not fixed constants), by Jensen's inequality:

$$\text{Var}(\mathbb{E}[X_j | \text{Pa}(j)] | S) > 0 \quad (39)$$

Rigorous justification:

For any $k \in M$, since $k \in \text{Pa}(j)$ and $\beta_{kj} \neq 0$ by assumption, the variable X_k appears in the structural equation for X_j with non-zero coefficient.

Define $Z := \sum_{k \in M} \beta_{kj} X_k$. We need to show that $\text{Var}(Z | S) > 0$.

If $\text{Var}(Z | S) = 0$, then Z would be almost surely constant given S . But this would imply that all X_k for $k \in M$ are deterministic functions of S (since the $\beta_{kj} \neq 0$). This contradicts the fact that

$M \subseteq \text{Pa}(j)$ are true causal parents with their own noise terms ϵ_k that are independent of S by the causal structure.

More formally: Each X_k for $k \in M$ has the form:

$$X_k = \exp \left(\theta_k + \sum_{\ell \in \text{Pa}(k)} \beta_{\ell k} X_\ell + \epsilon_k \right) \quad (40)$$

The noise ϵ_k is independent of all variables not in $\text{Desc}(k)$. Since $j \in \text{Desc}(k)$ (as $k \in \text{Pa}(j)$) but S may not fully contain $\text{Desc}(k)$, the conditional distribution of X_k given S retains variability from ϵ_k .

Therefore, $\text{Var}(Z|S) > 0$, which implies:

$$\text{Var}(\mathbb{E}[X_j|\text{Pa}(j)]|S) \geq \text{Var} \left(c_j \exp \left(\theta_j + \sum_{k \in S_{\text{in}}} \beta_{kj} X_k \right) \exp(Z) \middle| S \right) \quad (41)$$

$$= \left(c_j \exp \left(\theta_j + \sum_{k \in S_{\text{in}}} \beta_{kj} X_k \right) \right)^2 \cdot \text{Var}(\exp(Z)|S) \quad (42)$$

$$> 0 \quad (43)$$

where the strict inequality follows because $\text{Var}(\exp(Z)|S) > 0$ when $\text{Var}(Z|S) > 0$ (exponential is a non-constant monotone function).

Case 2: S contains no parents ($S \cap \text{Pa}(j) = \emptyset$)

The argument is similar but simpler. The entire term $\sum_{k \in \text{Pa}(j)} \beta_{kj} X_k$ is random given S , and by the same logic, this induces strictly positive conditional variance in $\mathbb{E}[X_j|\text{Pa}(j)]$ when conditioned on S .

Combining the pieces:

$$\mathbb{E}[\text{Var}(X_j|S)] = \mathbb{E}[\text{Var}(X_j|\text{Pa}(j))] + \mathbb{E}[\text{Var}(\mathbb{E}[X_j|\text{Pa}(j)]|S)] \quad (44)$$

We've shown that when $M \neq \emptyset$:

$$\mathbb{E}[\text{Var}(\mathbb{E}[X_j|\text{Pa}(j)]|S)] > 0 \quad (45)$$

Therefore:

$$\mathbb{E}[\text{Var}(X_j|S)] > \mathbb{E}[\text{Var}(X_j|\text{Pa}(j))] \quad (46)$$

Moreover, from part (ii), we have:

$$\mathbb{E}[(\mathbb{E}[X_j|S])^2] \leq \mathbb{E}[(\mathbb{E}[X_j|\text{Pa}(j)])^2] \quad (47)$$

Thus,

$$\mathcal{M}(j, S) = 1 + \frac{\mathbb{E}[\text{Var}(X_j|S)]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]} \quad (48)$$

$$> 1 + \frac{\mathbb{E}[\text{Var}(X_j|\text{Pa}(j))]}{\mathbb{E}[(\mathbb{E}[X_j|\text{Pa}(j)])^2]} \quad (49)$$

$$= \mathcal{M}(j, \text{Pa}(j)) \quad (50)$$

where the strict inequality follows from:

$$\frac{\mathbb{E}[\text{Var}(X_j|S)]}{\mathbb{E}[(\mathbb{E}[X_j|S])^2]} > \frac{\mathbb{E}[\text{Var}(X_j|\text{Pa}(j))]}{\mathbb{E}[(\mathbb{E}[X_j|\text{Pa}(j)])^2]} \quad (51)$$

because the numerator is strictly larger and the denominator is weakly smaller.

This completes the proof. $\square$

Appendix C. Proof of Proposition 2

Lemma (Ridge Consistency)

Under Assumption 2(B1)-(B2) and the log-linear model (2), the Ridge estimator $(\hat{\theta}_j, \hat{\beta}_j)$ in (4) with penalty $\lambda_{\text{ridge}} = o(1)$ as $n \rightarrow \infty$ satisfies $\|\hat{\beta}_j - \beta_j^*\|_2 = O_p(n^{-1/2})$ and $|\hat{\theta}_j - \theta_j^*| = O_p(n^{-1/2})$.

Proof. This follows from standard arguments for regularized least-squares regression on the log-scale (see e.g., [Bühlmann and van de Geer \(2011\)](#)). The sub-Gaussian design (B2) ensures bounded effective rank, while sub-exponential responses (B1) provide moment bounds. The rate $O_p(n^{-1/2})$ is standard for Ridge with vanishing penalty. $\square$

Proposition 2 (Concentration of empirical moment ratios). Consider the log-linear structural equation model (2) and fix a node j and candidate set S . Let $\hat{\mathcal{M}}(j, S)$ be the empirical moment-ratio estimator defined in (7), and $\mathcal{M}(j, S)$ its population counterpart in (6). Suppose Assumption 2 holds and the denominator $\mathbb{E}[(\mathbb{E}[X_j | S])^2]$ is bounded away from zero. Then there exist constants $C_{j,S}, c > 0$ such that, for all $\delta \in (0, 1)$ and all sufficiently large n ,

$$\mathbb{P}\left(|\hat{\mathcal{M}}(j, S) - \mathcal{M}(j, S)| \leq C_{j,S} \sqrt{\frac{\log(1/\delta)}{n}}\right) \geq 1 - \delta. \quad (52)$$

Proof We decompose the proof into five main steps.

Step 1: Notation and Decomposition.

Define:

- $A := \mathbb{E}[X_j^2]$ and $\hat{A}_n := \frac{1}{n} \sum_{i=1}^n (X_j^{(i)})^2$
- $\mu_{j|S}^{(i)} := \mathbb{E}[X_j | X_S^{(i)}]$ (the true conditional mean at sample i)
- $\hat{\mu}_{j|S}^{(i)} := \exp(\hat{\theta}_j + \sum_{k \in S} \hat{\beta}_{kj} X_k^{(i)})$ (the Ridge prediction)
- $B := \mathbb{E}[(\mu_{j|S})^2]$ and $\hat{B}_n := \frac{1}{n} \sum_{i=1}^n (\hat{\mu}_{j|S}^{(i)})^2$

Then:

$$\hat{\mathcal{M}}(j, S) = \frac{\hat{A}_n}{\hat{B}_n}, \quad \mathcal{M}(j, S) = \frac{A}{B} \quad (53)$$

We have:

$$\begin{aligned} \hat{\mathcal{M}}(j, S) - \mathcal{M}(j, S) &= \frac{\hat{A}_n}{\hat{B}_n} - \frac{A}{B} \\ &= \frac{B\hat{A}_n - A\hat{B}_n}{B\hat{B}_n} \\ &= \frac{B(\hat{A}_n - A) - A(\hat{B}_n - B)}{B\hat{B}_n} \end{aligned} \quad (54)$$

Step 2: Concentration of the Numerator $\hat{A}_n$.

By Assumption 2(B1), X_j is sub-exponential with parameters (ν_j, b_j) . This means that for all $t \geq 0$:

$$\mathbb{E}[\exp(t|X_j - \mathbb{E}[X_j]|)] \leq \exp(\nu_j^2 t^2 / 2) \quad \text{for all } t \in [0, 1/b_j] \quad (55)$$

Since X_j is sub-exponential, X_j^2 is also sub-exponential. Specifically, if $Y = X_j^2$, then Y is sub-exponential with parameters $(2\nu_j^2, 2b_j)$.

Applying Bernstein's inequality for sub-exponential random variables (?), we have:

$$\mathbb{P}\left(|\hat{A}_n - A| \geq t\right) \leq 2 \exp\left(-c_1 n \min\left(\frac{t^2}{\nu_j^4}, \frac{t}{b_j}\right)\right) \quad (56)$$

Setting the right-hand side equal to $\delta/3$ and solving for t :

$$t = C_1 \sqrt{\frac{\log(1/\delta)}{n}} \quad (57)$$

where C_1 depends on (ν_j, b_j) .

Therefore, with probability at least $1 - \delta/3$:

$$|\hat{A}_n - A| \leq C_1 \sqrt{\frac{\log(1/\delta)}{n}} \quad (58)$$

Step 3: Concentration of the Denominator $\hat{B}_n$.

We decompose:

$$\begin{aligned} \hat{B}_n - B &= \frac{1}{n} \sum_{i=1}^n (\hat{\mu}_{j|S}^{(i)})^2 - \mathbb{E}[(\mu_{j|S})^2] \\ &= \frac{1}{n} \sum_{i=1}^n \left[(\hat{\mu}_{j|S}^{(i)})^2 - (\mu_{j|S}^{(i)})^2 \right] + \frac{1}{n} \sum_{i=1}^n (\mu_{j|S}^{(i)})^2 - B \\ &=: T_1 + T_2 \end{aligned} \quad (59)$$

Step 3a: Bounding T_2 .

By Assumption 2(B1), $\mu_{j|S}$ is sub-exponential with parameters $(\nu_{j,S}, b_{j,S})$. By the same argument as for $\hat{A}_n$, $(\mu_{j|S})^2$ is sub-exponential, and applying Bernstein's inequality:

$$|T_2| \leq C_2 \sqrt{\frac{\log(1/\delta)}{n}} \quad \text{with probability } 1 - \delta/6 \quad (60)$$

Step 3b: Bounding T_1 (Prediction Error Term).

From the log-linear model:

$$\mu_{j|S}^{(i)} = \exp\left(\theta_j^* + \sum_{k \in S} \beta_{kj}^* X_k^{(i)}\right) \cdot \mathbb{E}[\exp(\epsilon_j)] \quad (61)$$

where (θ_j^*, β_j^*) are the population Ridge parameters.

Define the prediction error:

$$e_i := \hat{\mu}_{j|S}^{(i)} - \mu_{j|S}^{(i)} \quad (62)$$

We can write:

$$\begin{aligned} \hat{\mu}_{j|S}^{(i)} &= \exp \left(\hat{\theta}_j + \sum_{k \in S} \hat{\beta}_{kj} X_k^{(i)} \right) \\ &= \exp \left(\theta_j^* + \sum_{k \in S} \beta_{kj}^* X_k^{(i)} \right) \cdot \exp \left((\hat{\theta}_j - \theta_j^*) + \sum_{k \in S} (\hat{\beta}_{kj} - \beta_{kj}^*) X_k^{(i)} \right) \end{aligned} \quad (63)$$

Using the mean value theorem for the exponential function, for some ξ between the true and estimated log-scale predictions:

$$e_i = \mu_{j|S}^{(i)} \cdot \exp(\xi) \cdot \left[(\hat{\theta}_j - \theta_j^*) + \sum_{k \in S} (\hat{\beta}_{kj} - \beta_{kj}^*) X_k^{(i)} \right] \quad (64)$$

By Assumption 2(B2), the regressors X_S are sub-Gaussian, so $\|X_S^{(i)}\|_2 = O_p(\sqrt{|S|})$. By Cauchy-Schwarz:

$$\left| \sum_{k \in S} (\hat{\beta}_{kj} - \beta_{kj}^*) X_k^{(i)} \right| \leq \|\hat{\beta}_j - \beta_j^*\|_2 \cdot \|X_S^{(i)}\|_2 \quad (65)$$

By Lemma:

$$\|\hat{\beta}_j - \beta_j^*\|_2 = O_p(n^{-1/2}), \quad |\hat{\theta}_j - \theta_j^*| = O_p(n^{-1/2}) \quad (66)$$

Since $\mu_{j|S}^{(i)}$ is sub-exponential (B1), it has bounded exponential moments. Specifically, there exists $M_{j,S} > 0$ such that with high probability:

$$\mu_{j|S}^{(i)} \leq M_{j,S} \quad (67)$$

Combining these, for each i :

$$|e_i| \leq M_{j,S} \cdot \exp(|\xi|) \cdot O_p(n^{-1/2}) \cdot O_p(\sqrt{|S|}) = O_p(n^{-1/2}) \quad (68)$$

uniformly over i with high probability.

Now we bound T_1 :

$$\begin{aligned} T_1 &= \frac{1}{n} \sum_{i=1}^n \left[(\hat{\mu}_{j|S}^{(i)})^2 - (\mu_{j|S}^{(i)})^2 \right] \\ &= \frac{1}{n} \sum_{i=1}^n (\hat{\mu}_{j|S}^{(i)} + \mu_{j|S}^{(i)}) (\hat{\mu}_{j|S}^{(i)} - \mu_{j|S}^{(i)}) \\ &= \frac{1}{n} \sum_{i=1}^n (\hat{\mu}_{j|S}^{(i)} + \mu_{j|S}^{(i)}) \cdot e_i \end{aligned} \quad (69)$$

Using Cauchy-Schwarz:

$$\begin{aligned}
|T_1| &\leq \frac{1}{n} \sum_{i=1}^n |\hat{\mu}_{j|S}^{(i)} + \mu_{j|S}^{(i)}| \cdot |e_i| \\
&\leq \left(\frac{1}{n} \sum_{i=1}^n (\hat{\mu}_{j|S}^{(i)} + \mu_{j|S}^{(i)})^2 \right)^{1/2} \left(\frac{1}{n} \sum_{i=1}^n e_i^2 \right)^{1/2}
\end{aligned} \tag{70}$$

The first term is $O_p(1)$ (bounded empirical second moment by sub-exponential assumption), and the second term is $O_p(n^{-1/2})$ (since $e_i = O_p(n^{-1/2})$ uniformly). Therefore:

$$|T_1| = O_p(n^{-1/2}) \tag{71}$$

This is of smaller order than $\sqrt{\log(1/\delta)/n}$ for fixed δ , so for sufficiently large n :

$$|T_1| \leq C_3 \sqrt{\frac{\log(1/\delta)}{n}} \quad \text{with probability } 1 - \delta/6 \tag{72}$$

Step 3c: Combining for $\hat{B}_n$.

From (59), (60), and (72):

$$|\hat{B}_n - B| \leq |T_1| + |T_2| \leq (C_2 + C_3) \sqrt{\frac{\log(1/\delta)}{n}} \quad \text{with probability } 1 - \delta/3 \tag{73}$$

Step 4: Lower Bound on $\hat{B}_n$.

By assumption, $B = \mathbb{E}[(\mu_{j|S})^2]$ is bounded away from zero: $B \geq b_{\min} > 0$.

From (73), with probability $1 - \delta/3$:

$$\hat{B}_n \geq B - (C_2 + C_3) \sqrt{\frac{\log(1/\delta)}{n}} \geq \frac{B}{2} \geq \frac{b_{\min}}{2} \tag{74}$$

for sufficiently large n (specifically, $n \geq \frac{4(C_2+C_3)^2 \log(1/\delta)}{B^2}$).

Step 5: Final Bound on the Ratio.

From (54):

$$\begin{aligned}
|\hat{\mathcal{M}}(j, S) - \mathcal{M}(j, S)| &= \left| \frac{B(\hat{A}_n - A) - A(\hat{B}_n - B)}{B\hat{B}_n} \right| \\
&\leq \frac{|B(\hat{A}_n - A)| + |A(\hat{B}_n - B)|}{|B\hat{B}_n|} \\
&\leq \frac{B \cdot C_1 \sqrt{\frac{\log(1/\delta)}{n}} + A \cdot (C_2 + C_3) \sqrt{\frac{\log(1/\delta)}{n}}}{B \cdot \frac{B}{2}} \\
&= \frac{2C_1}{B} \sqrt{\frac{\log(1/\delta)}{n}} + \frac{2A(C_2 + C_3)}{B^2} \sqrt{\frac{\log(1/\delta)}{n}} \\
&\leq C_{j,S} \sqrt{\frac{\log(1/\delta)}{n}}
\end{aligned} \tag{75}$$

where $C_{j,S} := \frac{2C_1}{B} + \frac{2A(C_2+C_3)}{B^2}$ depends on the sub-exponential parameters and the moments of X_j and $\mu_{j|S}$.

By the union bound over the three probability events (Steps 2, 3a, 3b), this holds with probability at least:

$$1 - \left(\frac{\delta}{3} + \frac{\delta}{6} + \frac{\delta}{6} \right) = 1 - \frac{2\delta}{3} \quad (76)$$

Adjusting the constants appropriately, we can ensure the total failure probability is at most δ , completing the proof. $\blacksquare$

Appendix D. Proof of Proposition 3

H-MRS has two computational phases:

Phase 1 (Ordering): At each of p ordering steps m , we evaluate up to $p - m + 1$ remaining candidates. For each candidate j , we fit Ridge regression with at most $m - 1$ predictors to compute moment ratios. The total number of Ridge fits is:

$$\sum_{m=1}^p (p - m + 1) \leq p^2 \quad (77)$$

Each Ridge fit requires $O(nq^2 + q^3)$ operations, yielding $O(p^2 \cdot T_{\text{Ridge}})$ for Phase 1.

Phase 2 (Parent Selection): After ordering is complete, we fit ElasticNet once per variable to select parents from predecessors in the ordering. This requires p ElasticNet fits, each with cost $O(nq^2 \cdot K)$ using coordinate descent, yielding $O(p \cdot T_{\text{ElasticNet}})$ for Phase 2.

The total complexity is $O(p^2 \cdot T_{\text{Ridge}} + p \cdot T_{\text{ElasticNet}})$. Since ElasticNet typically requires more iterations than Ridge's closed-form solution, the ElasticNet term may dominate in practice, but the Ridge term scales quadratically with p . $\square$

Appendix E. Hyperparameter Selection Guidelines

This section provides detailed guidance on selecting hyperparameters for H-MRS, including both theoretical requirements for consistency and practical strategies for finite-sample settings.

The H-MRS algorithm involves several hyperparameters whose selection affects both practical performance and theoretical consistency. We now discuss the role of each parameter and provide guidance for their setting.

Ridge Regularization Parameter (λ_{ridge}). The Ridge parameter controls bias-variance tradeoff when estimating conditional expectations. For consistency of moment-ratio ordering:

- **Theoretical requirement:** $\lambda_{\text{ridge}} \rightarrow 0$ and $\lambda_{\text{ridge}} \cdot \sqrt{n} \rightarrow \infty$ as $n \rightarrow \infty$ ensures that Ridge estimates converge to true conditional expectations while maintaining stability.
- **Practical guidance:** Cross-validation on log-scale prediction error. In our experiments, $\lambda_{\text{ridge}} \in [0.01, 1.0]$ performed well.
- **Impact on identifiability:** Since moment-ratio minimization (Proposition 1) depends only on the ranking of $\mathcal{M}(j, S)$ values, moderate Ridge bias does not affect ordering consistency as long as the ranking is preserved.

ElasticNet Parameters (λ, ρ). These control sparsity in parent selection:

- **Theoretical requirement:** For consistency, $\lambda = \lambda_n$ should satisfy $\lambda_n \rightarrow 0$ and $\lambda_n \sqrt{n} \rightarrow \infty$, ensuring both convergence to true parameters and correct variable selection under standard ElasticNet theory (Zou and Hastie, 2005).
- **Practical guidance:** The mixing parameter $\rho \in [0.5, 0.9]$ balances grouping of correlated parents (low ρ) and strict sparsity (high ρ).
- **Impact on identifiability:** The plateau property (Proposition 1(iii)) guarantees that all supersets $S \supseteq \text{Pa}(j)$ achieve the same moment ratio. ElasticNet’s role is to select the minimal such set, which is a model selection problem separate from the identifiability established by moment ratios.

Thresholding Parameter (τ) and **Maximum Degree** ($d_{\max}$). These enforce sparsity in the final graph:

- **Theoretical requirement:** τ should exceed the ElasticNet estimation error: $\tau > C \cdot \sqrt{\frac{\log p}{n}}$ for some constant C ensures correct parent selection with high probability under standard sparsity assumptions.
- **Practical guidance:** Set τ as a small fraction of median non-zero ElasticNet coefficients (e.g., $\tau = 0.1 \cdot \text{median}(|\hat{\beta}_k| : \hat{\beta}_k \neq 0)$). The degree constraint $d_{\max}$ should reflect prior knowledge about graph sparsity.
- **Impact on identifiability:** These parameters do not affect the theoretical identifiability of causal ordering (Proposition 1) but control finite-sample edge recovery.

Summary of Consistency Conditions. For asymptotic consistency of H-MRS, hyperparameters should satisfy:

$$\lambda_{\text{ridge}}, \lambda_n \rightarrow 0, \quad \lambda_{\text{ridge}} \sqrt{n}, \lambda_n \sqrt{n} \rightarrow \infty, \quad \tau \gtrsim \sqrt{\frac{\log p}{n}} \quad (78)$$

These conditions ensure: (i) Ridge and ElasticNet estimators converge to population parameters, (ii) moment ratios concentrate around their population values (Proposition 2), and (iii) sparsity selection is consistent.

Appendix F. Estimated DAG for Financial Data Application




Figure 1: Causal structure discovered by H-MRS on financial data ($n=2,223$ companies, $p=19$ variables, 35 edges). The algorithm identified Equity Capital and Interest Expense as highly influential nodes with 13 and 15 outgoing edges respectively. Node colors indicate variable categories: blue=Assets, red=Liabilities, green=Equity, orange=Income, purple=Expenses, teal=Valuation.